

```
import numpy as np
import random
import gym
from gym.wrappers import Monitor
import glob
import io
import matplotlib.pyplot as plt
from IPython.display import HTML

def plot_Q(Q, message = "Q plot"):

    plt.figure(figsize=(10,10))
    plt.title(message)
    plt.pcolor(Q.max(-1), edgecolors='k', linewidths=2)
    plt.colorbar()
    # def x_direct(a):
    #     if a in [0, 2, 4, 5]:
    #         return 0
    #     return 1 if a == 1 else -1
    # def y_direct(a):
    #     if a in [1, 3]:
    #         return 0
    #     return 1 if (a == 0 or a == 4) else -1
    # policy = Q.argmax(-1)
    # policyx = np.vectorize(x_direct)(policy)
    # policyy = np.vectorize(y_direct)(policy)
    # idx = np.indices(policy.shape)
    # plt.quiver(idx[1].ravel()+0.5, idx[0].ravel()+0.5, policyx.ravel(), policyy.ravel(), pivot="middle", color='red')
    # plt.show()

def plot_visits(visits, message = "Average visits to each state"):
    plt.figure(figsize=(10,10))
    plt.title(message)
    plt.pcolor(visits, edgecolors='k', linewidths=2)
    plt.colorbar()
    plt.show()

env=gym.make('Taxi-v3')
env.reset()
state_shape = env.observation_space.n
action_shape = env.action_space.n
no_of_actions = env.action_space.n

#Current State
print(env.s)
print(env.action_space)

# 4x12 grid = 48 states
print ("Number of states:", env.nS)

# Primitive Actions
action = ["up", "right", "down", "left"]
#correspond to [0,1,2,3] that's actually passed to the environment

# either go left, up, down or right
print ("Number of actions that an agent can take:", env.nA)

# Example Transitions
rnd_action = random.randint(0, 3)
print ("Action taken:", action[rnd_action])
next_state, reward, is_terminal, t_prob = env.step(rnd_action)
print ("Transition probability:", t_prob)
print ("Next state:", next_state)
print ("Reward recieved:", reward)
print ("Terminal state:", is_terminal)
env.render()

267
Discrete(6)
Number of states: 500
Number of actions that an agent can take: 6
Action taken: right
Transition probability: {'prob': 1.0}
Next state: 167
Reward recieved: -1
Terminal state: False
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(North)

# We are defining options here
# Option 1 ["go_red"] - > Go to RED
# Option 2 ["go_yellow"] - > Go to YELLOW
# Option 3 ["go_green"] - > Go to GREEN
# Option 4 ["go_blue"] - > Go to BLUE

def go_red(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[1, 3, 0, 3, 3], [1, 3, 0, 3, 3], [1, 3, 3, 3, 3], [1, 1, 3, 1, 3], [1, 1, 3, 1, 3]]
    if r == 0 and c == 0:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

def go_yellow(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[0, 3, 0, 0, 0], [0, 3, 0, 3, 3], [0, 3, 3, 3, 3], [0, 1, 3, 1, 3], [0, 1, 3, 1, 3]]
    if r == 4 and c == 0:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

def go_green(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
```

```

actions = [[2, 0, 2, 2, 2], [2, 0, 2, 2, 1], [2, 2, 2, 2, 1], [1, 2, 1, 2, 1], [1, 2, 1, 2, 1]]
if r == 0 and c == 4:
    optdone = True
optact = actions[r][c]
return [optact, optdone]

def go_blue(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[2, 0, 2, 0, 3], [2, 0, 2, 0, 3], [2, 2, 2, 0, 3], [1, 2, 1, 0, 3], [1, 2, 1, 0, 3]]
    if r == 4 and c == 3:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

```

```
'''
```

```

Now the new action space will contain
["RED", "GREEN", "YELLOW", "BLUE", "pick passenger up", "drop passenger off"]
Corresponding to [0, 1, 2, 3, 4, 5]
'''

```

```

        '\nNow the new action space will contain\n["RED", "GREEN", "YELLOW", "BLUE", "pick passenger up", "drop passenger off"]\nCorresponding to [0, 1, 2, 3, 4, 5]\n'

```

```

rg = np.random.RandomState(42)

```

```

# Epsilon-greedy action selection function
def egreedy_policy(Q, state, epsilon):
    if not Q[state].any() or rg.rand() < epsilon:
        return rg.choice(Q.shape[-1])
    else:
        return np.argmax(Q[state])

```

▼ SMDP Q-learning

```

def SMDP_learning(env, q_values_SMDP, q_updates_SMDP, gamma, alpha, epsilon, episodes):
    rewards_SMDP=[]
    for _ in range(episodes):
        state = env.reset()
        done = False
        r = 0
        # While episode is not over
        while not done:
            # Choose action
            action = egreedy_policy(q_values_SMDP, state, epsilon=0.1)

            # Checking if primitive action
            if action < 6:
                # Perform regular Q-Learning update for state-action pair
                n_state, rew, done, _ = env.step(action)
                q_values_SMDP[state, action] = q_values_SMDP[state, action] + alpha*(rew + gamma*np.max(q_values_SMDP[n_state, :]) - q_values_SMDP[state, action])
                q_updates_SMDP[state, action] += 1
                state = n_state
                r += rew
                continue
            # Checking if action chosen is an option
            rew_bar = 0
            if action == 6: # action => Red option
                curr_state = state
                opdone = False
                n = 1
                while (opdone == False):
                    # Think about what this function might do?
                    opact, opdone = go_red(env, state)
                    n_state, rew, done,_ = env.step(opact)

                    # Is this formulation right? What is this term?
                    rew_bar = rew_bar + (gamma**(n - 1))*rew
                    n += 1
                    state = n_state
                    r += rew
                # Complete SMDP Q-Learning Update
                q_values_SMDP[curr_state, action] = q_values_SMDP[curr_state, action] + alpha*(rew_bar + (gamma**n)*np.max(q_values_SMDP[state, :]) - q_values_SMDP[curr_state, action])
                q_updates_SMDP[curr_state, action] += 1
                # Remember SMDP Updates. When & What do you update?
                continue
            if action == 7: # action => Yellow option
                curr_state = state
                opdone = False
                n = 1
                while (opdone == False):
                    # Think about what this function might do?
                    opact, opdone = go_yellow(env, state)
                    n_state, rew, done,_ = env.step(opact)

                    # Is this formulation right? What is this term?
                    rew_bar = rew_bar + (gamma**(n - 1))*rew
                    n += 1
                    state = n_state
                    r += rew
                # Complete SMDP Q-Learning Update
                q_values_SMDP[curr_state, action] = q_values_SMDP[curr_state, action] + alpha*(rew_bar + (gamma**n)*np.max(q_values_SMDP[state, :]) - q_values_SMDP[curr_state, action])
                q_updates_SMDP[curr_state, action] += 1
                continue
            if action == 8: # action => Green option
                curr_state = state
                opdone = False
                n = 1
                while (opdone == False):
                    # Think about what this function might do?
                    opact, opdone = go_green(env, state)
                    n_state, rew, done,_ = env.step(opact)

                    # Is this formulation right? What is this term?
                    rew_bar = rew_bar + (gamma**(n - 1))*rew
                    n += 1
                    state = n_state
                    r += rew
                # Complete SMDP Q-Learning Update
                q_values_SMDP[curr_state, action] = q_values_SMDP[curr_state, action] + alpha*(rew_bar + (gamma**n)*np.max(q_values_SMDP[state, :]) - q_values_SMDP[curr_state, action])
                q_updates_SMDP[curr_state, action] += 1

```

```
        continue
    if action == 9: # action => Blue option
        curr_state = state
        opdone = False
        n = 1
        while (opdone == False):
            # Think about what this function might do?
            opact, opdone = go_blue(env, state)
            n_state, rew, done,_ = env.step(opact)

            # Is this formulation right? What is this term?
            rew_bar = rew_bar + (gamma*(n - 1))*rew
            n += 1
            state = n_state
            r += rew
        # Complete SMDP Q-Learning Update
        q_values_SMDP[curr_state, action] = q_values_SMDP[curr_state, action] + alpha*(rew_bar + (gamma**n)*np.max(q_values_SMDP[state, :]) - q_values_SMDP[curr_state, action])
        q_updates_SMDP[curr_state, action] += 1
        continue
    rewards_SMDP.append(r)
return q_values_SMDP, q_updates_SMDP, rewards_SMDP
```


```
gamma = 0.90
alpha = 0.9
epsilon = 0.1
episodes = 2000
n_runs = 10
```

```
rewards_SMDP = []
Q_SMDP = []
Q_updates_SMDP = []
```

```
for n in range(n_runs):
    q_values_SMDP = np.zeros((state_shape, action_shape))
    q_updates_SMDP = np.zeros((state_shape, action_shape))
    q_values_SMDP, q_updates_SMDP, rewards = SMDP_learning(env, q_values_SMDP, q_updates_SMDP, gamma, alpha, epsilon, episodes)
    Q_SMDP.append(q_values_SMDP)
    Q_updates_SMDP.append(q_updates_SMDP)
    rewards_SMDP.append(rewards)
```

```
plt.plot(np.average(rewards_SMDP, 0))
plt.xlabel('Episodes')
plt.ylabel('Total Rewards')
plt.title('Reward Curve - Intra Option Q-learning')
```

Text(0.5, 1.0, 'Reward Curve - Intra Option Q-learning')




```
avg_reward_list_SMDP = []
idx = -1
reward_list = np.average(rewards_SMDP, 0)
for i in range(len(reward_list) - 99):
    temp = np.mean(reward_list[i:i + 100])
    avg_reward_list_SMDP.append(temp)
```


Plot of total reward vs episode

```
plt.plot(avg_reward_list_SMDP)
plt.xlabel('Episodes')
plt.ylabel('Total Reward')
plt.title('Running average of previous 100 rewards')
```

Text(0.5, 1.0, 'Running average of previous 100 rewards')



```
plot_Q(Q_SMDP[-1].reshape((25, 20, 6)))
```



▼ Another Set of Options

```
# We are defining options here
# Option 1 ["go_red2"] - > Go to RED
# Option 2 ["go_yellow2"] - > Go to YELLOW
# Option 3 ["go_green2"] - > Go to GREEN
# Option 4 ["go_blue2"] - > Go to BLUE

def go_red2(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[1, 3, 0, 0, 0], [1, 3, 0, 0, 0], [1, 3, 3, 3, 3], [1, 1, 1, 1, 1], [1, 1, 1, 1, 1]]
    if r == 0 and c == 0:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

def go_yellow2(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[0, 0, 0, 0, 0], [0, 0, 0, 0, 0], [0, 3, 3, 3, 3], [0, 1, 1, 1, 1], [0, 1, 1, 1, 1]]
    if r == 4 and c == 0:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

def go_green2(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[0, 0, 2, 2, 2], [0, 0, 2, 2, 1], [2, 2, 2, 2, 1], [1, 1, 1, 1, 1], [1, 1, 1, 1, 1]]
    if r == 0 and c == 4:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

def go_blue2(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[0, 0, 0, 0, 0], [0, 0, 0, 0, 0], [2, 2, 2, 0, 3], [1, 1, 1, 0, 3], [1, 1, 1, 0, 3]]
    if r == 4 and c == 3:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

'''
Now the new action space will contain
["RED", "GREEN", "YELLOW", "BLUE", "pick passenger up", "drop passenger off"]
Corresponding to [0, 1, 2, 3, 4, 5]
'''

'''
Now the new action space will contain\n["RED", "GREEN", "YELLOW", "BLUE", "pick passenger up", "drop passenger off"]\nCorr
esponding to [0 1 2 3 4 5]\n'''

def SMDP_learning2(env, q_values_SMDP, q_updates_SMDP, gamma, alpha, epsilon, episodes):
    rewards_SMDP=[]
    for _ in range(1000):
        state = env.reset()
        done = False
        r = 0
        # While episode is not over
        while not done:
            # Choose action
            action = egreedy_policy(q_values_SMDP, state, epsilon=0.1)

            # Checking if primitive action
            if action < 6:
                # Perform regular Q-Learning update for state-action pair
                n_state, rew, done, _ = env.step(action)
                q_values_SMDP[state, action] = q_values_SMDP[state, action] + alpha*(rew + gamma*np.max(q_values_SMDP[n_state, :]) - q_values_SMDP[state, action])
                q_updates_SMDP[state, action] += 1
                state = n_state
                r += rew
                continue
            # Checking if action chosen is an option
            rew_bar = 0
            if action == 6: # action => Red option
                curr_state = state
                optdone = False
```

```

opdone = raise
n = 1
while (opdone == False):
    # Think about what this function might do?
    opact, opdone = go_red2(env, state)
    n_state, rew, done,_ = env.step(opact)

    # Is this formulation right? What is this term?
    rew_bar = rew_bar + (gamma**(n - 1))*rew
    n += 1
    state = n_state
    r += rew
# Complete SMDP Q-Learning Update
q_values_SMDP[curr_state, action] = q_values_SMDP[curr_state, action] + alpha*(rew_bar + (gamma**n)*np.max(q_values_SMDP[state, :]) - q_values_SMDP[curr_state, action])
q_updates_SMDP[curr_state, action] += 1
# Remember SMDP Updates. When & What do you update?
continue
if action == 7: # action => Yellow option
    curr_state = state
    opdone = False
    n = 1
    while (opdone == False):
        # Think about what this function might do?
        opact, opdone = go_yellow2(env, state)
        n_state, rew, done,_ = env.step(opact)

        # Is this formulation right? What is this term?
        rew_bar = rew_bar + (gamma**(n - 1))*rew
        n += 1
        state = n_state
        r += rew
    # Complete SMDP Q-Learning Update
    q_values_SMDP[curr_state, action] = q_values_SMDP[curr_state, action] + alpha*(rew_bar + (gamma**n)*np.max(q_values_SMDP[state, :]) - q_values_SMDP[curr_state, action])
    q_updates_SMDP[curr_state, action] += 1
    continue
if action == 8: # action => Green option
    curr_state = state
    opdone = False
    n = 1
    while (opdone == False):
        # Think about what this function might do?
        opact, opdone = go_green2(env, state)
        n_state, rew, done,_ = env.step(opact)

        # Is this formulation right? What is this term?
        rew_bar = rew_bar + (gamma**(n - 1))*rew
        n += 1
        state = n_state
        r += rew
    # Complete SMDP Q-Learning Update
    q_values_SMDP[curr_state, action] = q_values_SMDP[curr_state, action] + alpha*(rew_bar + (gamma**n)*np.max(q_values_SMDP[state, :]) - q_values_SMDP[curr_state, action])
    q_updates_SMDP[curr_state, action] += 1
    continue
if action == 9: # action => Blue option
    curr_state = state
    opdone = False
    n = 1
    while (opdone == False):
        # Think about what this function might do?
        opact, opdone = go_blue2(env, state)
        n_state, rew, done,_ = env.step(opact)

        # Is this formulation right? What is this term?
        rew_bar = rew_bar + (gamma**(n - 1))*rew
        n += 1
        state = n_state
        r += rew
    # Complete SMDP Q-Learning Update
    q_values_SMDP[curr_state, action] = q_values_SMDP[curr_state, action] + alpha*(rew_bar + (gamma**n)*np.max(q_values_SMDP[state, :]) - q_values_SMDP[curr_state, action])
    q_updates_SMDP[curr_state, action] += 1
    continue
rewards_SMDP.append(r)
return q_values_SMDP, q_updates_SMDP, rewards_SMDP

```

```

gamma = 0.90
alpha = 0.9
epsilon = 0.1
episodes = 2000
n_runs = 10

```

```

rewards_SMDP = []
Q_SMDP = []
Q_updates_SMDP = []

```

```

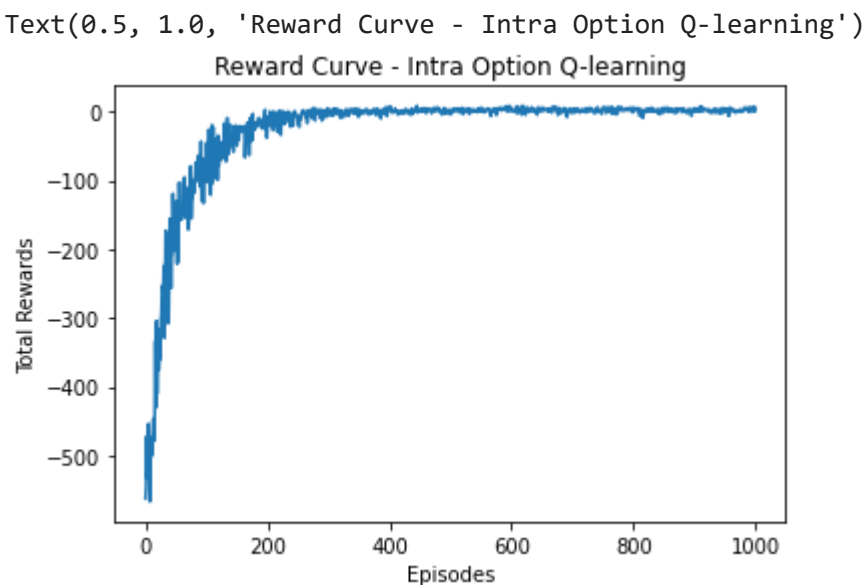
for n in range(n_runs):
    q_values_SMDP = np.zeros((state_shape, action_shape))
    q_updates_SMDP = np.zeros((state_shape, action_shape))
    q_values_SMDP, q_updates_SMDP, rewards = SMDP_learning2(env, q_values_SMDP, q_updates_SMDP, gamma, alpha, epsilon, episodes)
    Q_SMDP.append(q_values_SMDP)
    Q_updates_SMDP.append(q_updates_SMDP)
    rewards_SMDP.append(rewards)

```

```

plt.plot(np.average(rewards_SMDP, 0))
plt.xlabel('Episodes')
plt.ylabel('Total Rewards')
plt.title('Reward Curve - Intra Option Q-learning')

```

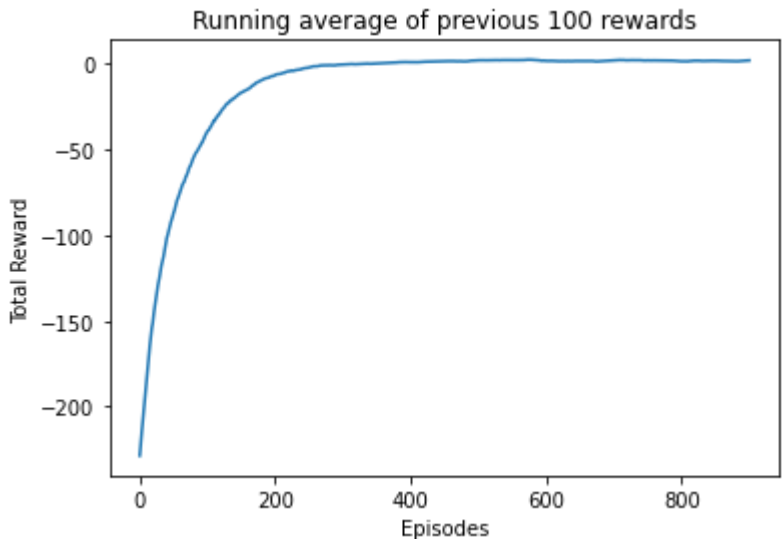


```
avg_reward_list_SMDP = []
idx = -1
reward_list = np.average(rewards_SMDP, 0)
for i in range(len(reward_list) - 99):
    temp = np.mean(reward_list[i:i + 100])
    avg_reward_list_SMDP.append(temp)
```


Plot of total reward vs episode

```
plt.plot(avg_reward_list_SMDP)
plt.xlabel('Episodes')
plt.ylabel('Total Reward')
plt.title('Running average of previous 100 rewards')
```

Text(0.5, 1.0, 'Running average of previous 100 rewards')



```
plot_Q(Q_SMDP[-1].reshape((25, 20, 6)))
```



▼ Intra-option Q-learning

```
import numpy as np
import random
import gym
from gym.wrappers import Monitor
import glob
import io
import matplotlib.pyplot as plt
from IPython.display import HTML

'''
The environment used here is extremely similar to the openai gym ones.
At first glance it might look slightly different.
The usual commands we use for our experiments are added to this cell to aid you
work using this environment.
'''

#Setting up the environment
env = gym.make('Taxi-v3')
env.seed(0)

state_shape = env.observation_space.n
action_shape = env.action_space.n
no_of_actions = env.action_space.n
"0: south; 1: north; 2: east; 3: west; 4: pick passenger up; and 5: drop passenger off"

print(state_shape)
print(no_of_actions)
print(env.action_space.sample())
print("----")

state = env.reset()
''' This returns the initial state (when environment is reset) '''

print(list(env.decode(state)))
print("----")

action = env.action_space.sample()

print(action)
print("----")
```



```
next_state, reward, done, info = env.step(action)
''' env.step is used to calculate new state and obtain reward based on old state and action taken '''

print(list(env.decode(next_state)))
print(reward)
print(done)
print(info)
print("----")

500
6
5
----
[0, 1, 1, 2]
----
1
----
[0, 1, 1, 2]
-1
False
{'prob': 1.0}
----

s = env.reset()
print(s)
env.render()

312
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

"0: south; 1: north; 2: east; 3: west; 4: pick passenger up; and 5: drop passenger off"
next_state, reward, done, info = env.step(3)
env.render()
print(next_state)
print(list(env.decode(next_state)))
print(reward)
print(done)

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(West)
312
[3, 0, 3, 0]
-1
False

▼ Options

We custom define very simple options here. They might not be the logical options for this settings deliberately chosen to visualise the Q Table better.

# We are defining options here
# Option 1 ["go_red"] - > Go to RED
# Option 2 ["go_yellow"] - > Go to YELLOW
# Option 3 ["go_green"] - > Go to GREEN
# Option 4 ["go_blue"] - > Go to BLUE

def go_red(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[1, 3, 0, 3, 3], [1, 3, 0, 3, 3], [1, 3, 3, 3, 3], [1, 1, 3, 1, 3], [1, 1, 3, 1, 3]]
    if r == 0 and c == 0:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

def go_yellow(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[0, 3, 0, 0, 0], [0, 3, 0, 3, 3], [0, 3, 3, 3, 3], [0, 1, 3, 1, 3], [0, 1, 3, 1, 3]]
    if r == 4 and c == 0:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

def go_green(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[2, 0, 2, 2, 2], [2, 0, 2, 2, 1], [2, 2, 2, 2, 1], [1, 2, 1, 2, 1], [1, 2, 1, 2, 1]]
    if r == 0 and c == 4:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

def go_blue(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[2, 0, 2, 0, 3], [2, 0, 2, 0, 3], [2, 2, 2, 0, 3], [1, 2, 1, 0, 3], [1, 2, 1, 0, 3]]
    if r == 4 and c == 3:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]

'''
Now the new action space will contain
["RED", "GREEN", "YELLOW", "BLUE", "pick passenger up", "drop passenger off"]
Corresponding to [0, 1, 2, 3, 4, 5]
'''

'\nNow the new action space will contain\n["RED", "GREEN", "YELLOW", "BLUE", "pick passenger up", "drop passenger off"]\nCorresponding to [0, 1, 2, 3, 4, 5]\n'
```

▼ Intra-option Q-learning

```
rg = np.random.RandomState(42)

# Epsilon-greedy action selection function
def egreedy_policy(Q, state, epsilon):
    if not Q[state].any() or rg.rand() < epsilon:
        return rg.choice(Q.shape[-1])
    else:
        return np.argmax(Q[state])

def IO_learning(env, q_values_IO, q_updates_IO, gamma, alpha, epsilon, episodes):
    rewards_IO = []
    # Iterate over episodes
    for ep in range(episodes):
        state = env.reset()
        done = False
        r = 0
        # While episode is not over
        while not done:
            action = egreedy_policy(q_values_IO, state, epsilon = epsilon)

            # Checking if primitive action
            if action > 3:
                next_state, reward, done, _ = env.step(action)
                q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
                q_updates_IO[state, action] += 1
                state = next_state
                r += reward
                continue
            # Checking if action chosen is an option
            if action == 0: # action => RED option
                optdone = False
                while (optdone == False):
                    optact, optdone = go_red(env, state)
                    next_state, reward, done, _ = env.step(optact)
                    r, c, _, _ = list(env.decode(next_state))
                    if r == 0 and c == 0:
                        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
                        optdone = True
                    else:
                        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*q_values_IO[next_state, action] - q_values_IO[state, action])
                        q_updates_IO[state, action] += 1
                        state = next_state
                        r += reward
                continue
            if action == 1:
                optdone = False
                while (optdone == False): # action => GREEN option
                    optact, optdone = go_green(env, state)
                    next_state, reward, done, _ = env.step(optact)
                    r, c, _, _ = list(env.decode(next_state))
                    if r == 0 and c == 4:
                        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
                        optdone = True
                    else:
                        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*q_values_IO[next_state, action] - q_values_IO[state, action])
                        q_updates_IO[state, action] += 1
                        state = next_state
                        r += reward
                continue
            if action == 2: # action => YELLOW option
                optdone = False
                while (optdone == False):
                    optact, optdone = go_yellow(env, state)
                    next_state, reward, done, _ = env.step(optact)
                    r, c, _, _ = list(env.decode(next_state))
                    if r == 4 and c == 0:
                        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
                        optdone = True
                    else:
                        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*q_values_IO[next_state, action] - q_values_IO[state, action])
                        q_updates_IO[state, action] += 1
                        state = next_state
                        r += reward
                continue
            if action == 3: # action => BLUE option
                optdone = False
                while (optdone == False):
                    optact, optdone = go_blue(env, state)
                    next_state, reward, done, _ = env.step(optact)
                    r, c, _, _ = list(env.decode(next_state))
                    if r == 4 and c == 3:
                        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
                        optdone = True
                    else:
                        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*q_values_IO[next_state, action] - q_values_IO[state, action])
                        q_updates_IO[state, action] += 1
                        state = next_state
                        r += reward
                continue
            rewards_IO.append(r)
        return q_values_IO, q_updates_IO, rewards_IO

#### Intra-Option Q-Learning

# Divergence when alpha > 1.2

# Hyperparameters
gamma = 0.90
alpha = 0.9
epsilon = 0.1
episodes = 2000
n_runs = 10


rewards_IO = []
Q_IO = []
Q_updates_IO = []
```



```
for n in range(n_runs):
    q_values_I0 = np.zeros((state_shape, action_shape))
    q_updates_I0 = np.zeros((state_shape, action_shape))
    q_values_I0, q_updates_I0, rewards = IO_learning(env, q_values_I0, q_updates_I0, gamma, alpha, epsilon, episodes)
    Q_I0.append(q_values_I0)
    Q_updates_I0.append(q_updates_I0)
    rewards_I0.append(rewards)
```

▼ Q-Tables and Update Frequencies

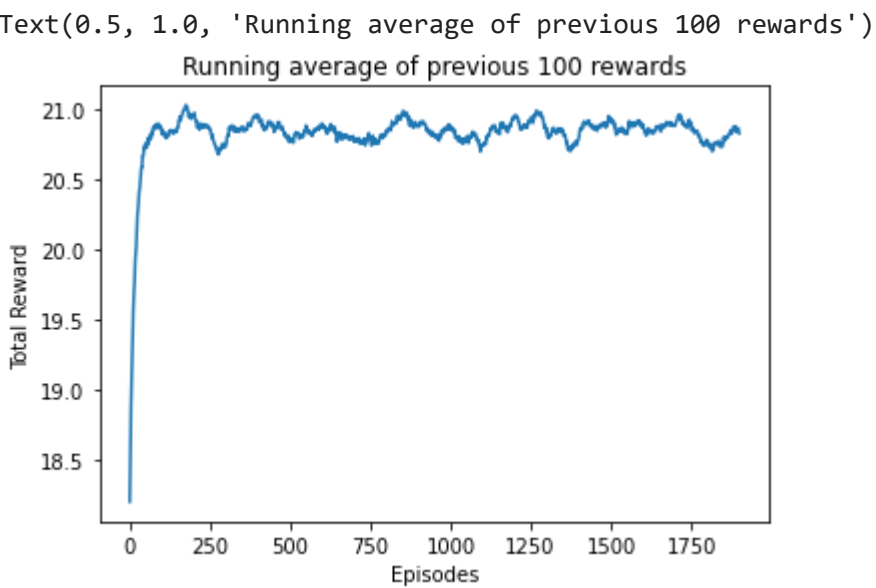
```
plt.plot(np.average(rewards_I0, 0))
plt.xlabel('Episodes')
plt.ylabel('Total Rewards')
plt.title('Reward Curve - Intra Option Q-learning')
```




```
avg_reward_list = []
idx = -1
reward_list = np.average(rewards_I0, 0)
for i in range(len(reward_list) - 99):
    temp = np.mean(reward_list[i:i + 100])
    avg_reward_list.append(temp)
```

Plot of total reward vs episode

```
plt.plot(avg_reward_list)
plt.xlabel('Episodes')
plt.ylabel('Total Reward')
plt.title('Running average of previous 100 rewards')
```



```
plot_Q(Q_I0[-1].reshape((25, 20, 6)))
```




```
list(env.decode(479)) # 16, 97, 418, 479
```

```
[4, 3, 4, 3]
```

```
Q_I0[-1][479]
```

```
array([-3.14409118, -0.1359871 , -3.14645781, 16.9821      ,  7.9999983 ,
        20.          ])
```

```
plot_visits(np.sum(Q_updates_I0[-1], axis = -1).reshape((25, 20)))
```



▼ Another set of Options

```
# We are defining options here
# Option 1 ["go_red2"] - > Go to RED
# Option 2 ["go_yellow2"] - > Go to YELLOW
# Option 3 ["go_green2"] - > Go to GREEN
# Option 4 ["go_blue2"] - > Go to BLUE
```

```
def go_red2(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[1, 3, 0, 0, 0], [1, 3, 0, 0, 0], [1, 3, 3, 3, 3], [1, 1, 1, 1, 1], [1, 1, 1, 1, 1]]
    if r == 0 and c == 0:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]
```

```
def go_yellow2(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[0, 0, 0, 0, 0], [0, 0, 0, 0, 0], [0, 3, 3, 3, 3], [0, 1, 1, 1, 1], [0, 1, 1, 1, 1]]
    if r == 4 and c == 0:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]
```

```
def go_green2(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[0, 0, 2, 2, 2], [0, 0, 2, 2, 1], [2, 2, 2, 2, 1], [1, 1, 1, 1, 1], [1, 1, 1, 1, 1]]
    if r == 0 and c == 4:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]
```

```
def go_blue2(env, state):
    r, c, _, _ = list(env.decode(state))
    optdone = False
    actions = [[0, 0, 0, 0, 0], [0, 0, 0, 0, 0], [2, 2, 2, 0, 3], [1, 1, 1, 0, 3], [1, 1, 1, 0, 3]]
    if r == 4 and c == 3:
        optdone = True
    optact = actions[r][c]
    return [optact, optdone]
```

```
...
```

Now the new action space will contain

["RED", "GREEN", "YELLOW", "BLUE", "pick passenger up", "drop passenger off"]

Corresponding to [0, 1, 2, 3, 4, 5]

```
...
```

```
'\nNow the new action space will contain\n["RED", "GREEN", "YELLOW", "BLUE", "pick passenger up", "drop passenger off"]\nCorresponding to [0, 1, 2, 3, 4, 5]\n'
```

```
def IO_learning2(env, q_values_IO, q_updates_IO, gamma, alpha, epsilon, episodes):
    rewards_IO = []
    # Iterate over episodes
    for ep in range(episodes):
        state = env.reset()
        done = False
        r = 0
        # While episode is not over
        while not done:
            action = egreedy_policy(q_values_IO, state, epsilon = epsilon)

            # Checking if primitive action
            if action > 3:
                next_state, reward, done, _ = env.step(action)
                q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
                q_updates_IO[state, action] += 1
                state = next_state
                r += reward
                continue
            # Checking if action chosen is an option
            if action == 0: # action => RED option
                optdone = False
                while (optdone == False):
                    optact, optdone = go_red2(env, state)
                    next_state, reward, done, _ = env.step(optact)
                    r, c, _, _ = list(env.decode(next_state))
                    if r == 0 and c == 0:
```

```

        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
        optdone = True
    else:
        q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*q_values_IO[next_state, action] - q_values_IO[state, action])
    q_updates_IO[state, action] += 1
    state = next_state
    r += reward
    continue
if action == 1:
    optdone = False
    while (optdone == False): # action => GREEN option
        optact, optdone = go_green2(env, state)
        next_state, reward, done,_ = env.step(optact)
        r, c, _, _ = list(env.decode(next_state))
        if r == 0 and c == 4:
            q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
            optdone = True
        else:
            q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*q_values_IO[next_state, action] - q_values_IO[state, action])
    q_updates_IO[state, action] += 1
    state = next_state
    r += reward
    continue
if action == 2: # action => YELLOW option
    optdone = False
    while (optdone == False):
        optact, optdone = go_yellow2(env, state)
        next_state, reward, done,_ = env.step(optact)
        r, c, _, _ = list(env.decode(next_state))
        if r == 4 and c == 0:
            q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
            optdone = True
        else:
            q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*q_values_IO[next_state, action] - q_values_IO[state, action])
    q_updates_IO[state, action] += 1
    state = next_state
    r += reward
    continue
if action == 3: # action => BLUE option
    optdone = False
    while (optdone == False):
        optact, optdone = go_blue2(env, state)
        next_state, reward, done,_ = env.step(optact)
        r, c, _, _ = list(env.decode(next_state))
        if r == 4 and c == 3:
            q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*np.max(q_values_IO[next_state, :]) - q_values_IO[state, action])
            optdone = True
        else:
            q_values_IO[state, action] = q_values_IO[state, action] + alpha*(reward + gamma*q_values_IO[next_state, action] - q_values_IO[state, action])
    q_updates_IO[state, action] += 1
    state = next_state
    r += reward
    continue
    rewards_IO.append(r)
return q_values_IO, q_updates_IO, rewards_IO

```

Intra-Option Q-Learning

```

# Hyperparameters
gamma = 0.90
alpha = 0.90
epsilon = 0.1
episodes = 2000
n_runs = 10

rewards_IO = []
Q_IO = []
Q_updates_IO = []

for n in range(n_runs):
    q_values_IO = np.zeros((state_shape, action_shape))
    q_updates_IO = np.zeros((state_shape, action_shape))
    q_values_IO, q_updates_IO, rewards = IO_learning2(env, q_values_IO, q_updates_IO, gamma, alpha, epsilon, episodes)
    Q_IO.append(q_values_IO)
    Q_updates_IO.append(q_updates_IO)
    rewards_IO.append(rewards)

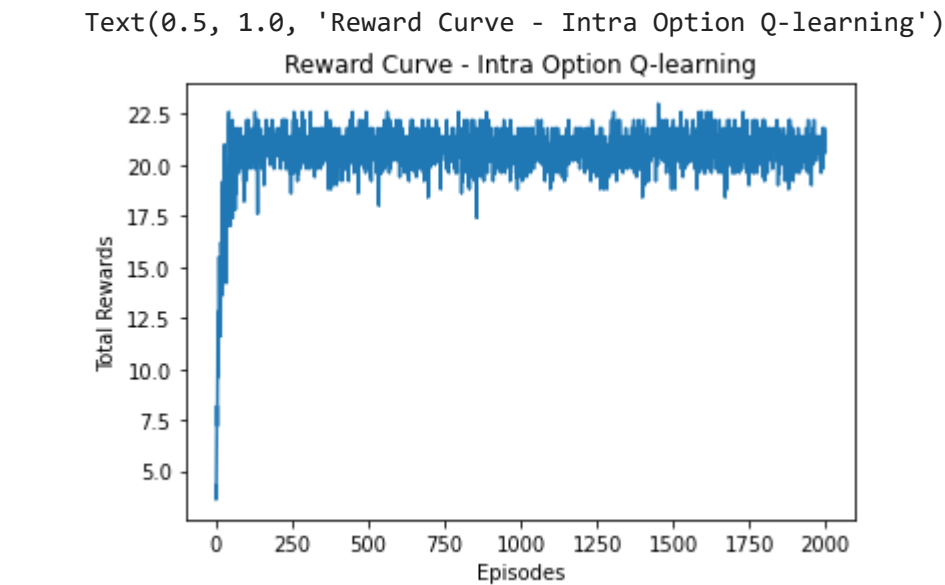
```

▼ Q-Tables and Update Frequencies

```

plt.plot(np.average(rewards_IO, 0))
plt.xlabel('Episodes')
plt.ylabel('Total Rewards')
plt.title('Reward Curve - Intra Option Q-learning')

```




```


avg_reward_list = []
idx = -1
reward_list = np.average(rewards_IO, 0)
for i in range(len(reward_list) - 99):
    temp = np.mean(reward_list[i:i + 100])
    avg_reward_list.append(temp)

```

Plot of total reward vs episode



```
plot_Q(Q_I0[-1].reshape((25, 20, 6)))
```



```
plot_visits(np.sum(Q_updates_I0[-1], axis = -1).reshape((25, 20)))
```

